The black box we choose to believe in

When AI fails in hospitals, courtrooms, and customer service, who is held responsible? TRUST AI? — PART 2 OF 4

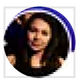

THE TRUST FRONTIER
JUN 16, 2026

A man named Jake Moffatt was on his way to his grandmother's funeral when he asked Air Canada's chatbot whether he could claim a bereavement discount after the fact. The bot said yes. The airline said no. When he took them to court, Air Canada argued the chatbot was a separate legal entity responsible for its own responses.

I read that sentence three times.

Not because it is legally absurd, though it is. Because it is honest. A company deployed a system it was not willing to stand behind — and when that system caused harm, its first instinct was to point at the machine [The tribunal rejected that argument](Moffatt v. Air Canada, 2024 BCCRT 149) and ordered Air Canada to pay $812 in total compensation, of which $650.88 was awarded specifically for the chatbot's negligent misrepresentation. The airline's legal fees were almost certainly higher.

I have been thinking about that case since I started researching Part 2 of this series. Part 1 asked why we trust AI so readily. The answer was design, architecture, and habit. This piece asks a harder question: what happens when that trust breaks, and the thing we trusted turns out to be pointing back at itself?

The Machine That Invented the Law

A senior research fellow at the Paris School of Advanced Business Studies (HEC Paris), lists more than 1,400 cases where courts have addressed AI errors in the past three years, including filings by attorneys and self-represented litigants. The pattern is almost always the same: a lawyer uses an AI tool to research case law, the tool invents citations that sound exactly like real ones, and the lawyer submits them without checking. The cases do not exist. The judges do not find them. The lawyer finds out in court.

In April 2026, [the Alabama Supreme Court dismissed an appeal and sanctioned attorney W. Perry Hall](Bloomberg Law) for filing briefs described as "grossly deficient" — containing what the court called an "astounding number" of fabricated legal authorities. The court ordered him to pay $17,200 in fees and referred him to the state bar. His clients lost their appeal not because their case lacked merit, but because the citations their lawyer submitted did not exist.

The detail that stays with me is in the concurring opinion.[In the same footnote where Hall apologized for citing AI-hallucinated authorities and promised "the mistake will not recur," he proceeded to cite two more cases that did not exist — in the very next sentence.](Yellowhammer News) Justice Greg Cook wrote: "It is simply hard to imagine how this could occur absent, perhaps, using AI to craft the apology for having used AI."

> *The machine invented the citations. Then it may have written the apology for inventing them.*

[This is not an isolated case of one careless lawyer. Courts have applied existing rules inconsistently, with no explicit federal guidance on how to address the misuse of generative AI in court filings.](Sterne Kessler) One federal judge put it plainly after nearly applying hallucinated citations in a court order herself: "Plaintiff's use of AI affirmatively misled me."

I keep returning to that. Not as a call for prosecution, but as a fact about accountability. When a human lawyer fabricates a citation, we know who is responsible. When an AI fabricates one and a human submits it unchecked, the question gets distributed across a chain of people and systems until nobody is clearly, legally, individually at fault.

The Scan That Missed It

The courtroom is a contained setting. When a citation is fabricated, someone can catch it — eventually. The hospital is different. In a hospital, the alert fires and a clinician has to decide, in real time, whether to trust it.

In Part 1 of this series, I wrote about performance-based trust: how we have shifted from evaluating AI on accountability to evaluating it on outputs, mistaking confident fluency for verified competence. I did not have a better illustration of what that costs than this.

[Epic, the nation's largest seller of medical records, developed an algorithm to predict which patients would develop sepsis — the leading cause of in-hospital deaths — so doctors could intervene earlier. Hundreds of hospitals plugged the algorithm in without verifying its advertised accuracy rate.](STAT News) Epic claimed the model was 76 to 83 percent accurate. There had been no credible independent tests of any of its algorithms — until a team at the University of Michigan started asking questions.

[What they found: the model missed two-thirds of sepsis cases. Its overall accuracy was about 63 percent — little better than a coin flip. And its false alarm rate was so high that clinicians would have to respond to 109 alerts to find a single patient who actually had sepsis.](Wong et al., JAMA Internal Medicine, 2021)

The reason matters as much as the number. Epic used billing codes to define what sepsis looked like — not the clinical definition that is meaningful to a patient or a health system. The model was not trained to detect early sepsis. It was trained to predict billing code decisions.

The hospitals that deployed it did not know that. They were not told to ask.

There is currently no harmonized regulatory structure governing who is responsible when an AI system integrated into clinical practice causes harm. When the scan misses the tumour, the question loops through the hospital that deployed the system, the company that built it, the doctor who trusted it, and the regulatory body that approved or failed to approve it. None of them is cleanly at fault. All of them are partly. And the patient is left navigating that loop while managing a diagnosis that arrived too late.

This is the part of the story that does not fit neatly into a conversation about AI error. The model did not malfunction. It performed exactly as designed. The problem was that what it was designed to do and what the hospitals believed it was doing were two different things entirely — and no one closed that gap before the alerts started firing.

The Pattern Beneath All Three Cases

Three cases. Three settings. And in each one, the same thing happened: the closer the AI failure got to a human body, the harder it became to name who was responsible. Air Canada paid $812. That is the cost of accountability when the stakes are low enough to fit inside a small claims tribunal. Hall's clients lost their appeal. The sepsis patients — we do not have a number for them, because the harm was never attributed. That is not a coincidence. It is a gradient.

The yes-man in the machine does not just tell you what you want to hear. It tells your lawyer what the law says. It tells your doctor what the scan shows. It tells the grieving passenger what the refund policy is. When it is wrong in those moments, the question of who is responsible loops back through a chain of developers, deployers, regulators, and users that nobody has fully mapped yet.

The law was built for a world where humans made decisions and other humans could be held responsible. AI broke that assumption quietly. The EU's AI Act classifies AI systems by risk and imposes obligations on developers and deployers of high-risk applications in healthcare, justice, and employment. It is the only major jurisdiction to have done so in binding law. The United States has no equivalent framework. What fills that gap is litigation — slow, expensive, and always arriving after the harm has already occurred. That is not a governance system. It is the absence of one.

This is not a technical problem. It is a trust problem. And unlike the dependency cycle described in Part 1 — which we did to ourselves gradually, through convenience and habit — this one was built into the system from the start. The accountability gap was not an accident. It was the default.

Have you encountered an AI failure in a professional or high-stakes setting? I am collecting cases for the next part of this series. Reply and tell me what happened.

Next in this series: "The ghost in the governance gap" — why the EU and US are building opposite AI systems, and what the data says about how much people in each jurisdiction actually trust the protection they have.

Sources

[1] Moffatt v. Air Canada, 2024 BCCRT 149 (British Columbia Civil Resolution Tribunal)

[2] Charlotin, D. AI Hallucination Cases Database, 2025 — 700+ cases identified globally

[3] Scientific American, “Why Lawyers Keep Citing Fake Cases Invented by AI,” May 2025

[4] Bloomberg Law, Alabama Supreme Court ruling, April 2026

[5] Yellowhammer News, Justice Cook concurring opinion, April 2026

[6] Sterne Kessler, AI Hallucinations in Court Filings: 2025 Review

[7] Wong et al., “External Validation of a Widely Implemented Proprietary Sepsis Prediction Model,” JAMA Internal Medicine, 2021

[8] STAT News, Epic sepsis algorithm investigation, 2021

[9] IHPI, University of Michigan sepsis model analysis

[10] NIH/NCBI, “Artificial intelligence in hospitals: Legal uncertainties and emerging risks,” July 2025

Disclosure: AI tools were used to assist in editing and formatting. All arguments, source selection, and analytical judgments are the author’s own, drawn from original research.