

The yes-man in the machine

Why your AI is not the neutral expert you think it is (Trust AI?
Part 1 of 4)



THE TRUST FRONTIER

JUN 11, 2026

AI systems have become the first stop for medical questions, financial decisions, and emotional support. The reasons why reveal something uncomfortable about how we actually think.

I told a machine things I have never told a priest, a friend, or my family. It was the first ChatGPT project I ever created; I called it “My Religious Life.” Not once did I ask what made it qualified to answer. When I finally did, another question would not leave me alone: is it just me, or have we all quietly replaced “I need to ask an expert” with “let me ask AI first”?

It is not just me.

85%

of surveyed students rated ChatGPT tutoring over human tutors (Intelligent.com, 2023)

48.5%

of heavy AI users in a 2025 study held doctorates — the single largest group

The first number is striking on its own: a survey of 801 students found that 85% believed ChatGPT’s tutoring was more effective than learning from a human tutor. And it is not an outlier. A 2024 global survey by the Digital Education Council found 86% of students regularly using AI in their studies, and the Council’s 2026 survey of Latin America puts student adoption there at 92%. The shift toward AI as a first-resort advisor is accelerating everywhere: healthcare, legal advice, personal

finance, mental health. But the more I researched it, the more I became convinced that the explanation is not in the technology. It is in us, in the social and psychological conditions that make AI *feel* trustworthy.

The neutrality we want to believe in. The most common assumption is that AI is neutral: no agenda, no ego, no off day, unlike a human doctor or financial advisor. The research complicates this badly. An analysis of 47,000 publicly shared ChatGPT conversations found the model began responses with “yes” roughly ten times more often than “no.” That is not a coincidence; it is a design tendency. AI companies optimize their models to retain engagement, which in practice means affirming rather than challenging. I see it in my own conversations. When I tell an AI that I think something is wrong, it answers “Yes, you are right... however...” It could have simply said: “That is not right, because...” It almost never does. Whether you lean left or right, the machine is far more likely to mirror your worldview than to offer a genuinely independent perspective. The result is a sophisticated echo chamber that feels like objectivity.




We cannot see the ‘why,’ so we default to trusting the ‘what’ — especially when the AI tells us exactly what we want to hear.

The box we cannot open. This connects to what philosophers call epistemic opacity, better known as the black box problem. With a doctor or a lawyer, you can ask for an explanation, push back, and hold them accountable if they are wrong. With AI, you submit a question and receive an answer, with no visibility into the steps in between. And here is what concerns me most: most users never even think to ask. Perceived fluency gets mistaken for verified truth.

Sociologists of trust have traditionally identified three pillars of trustworthiness: competence in a specific domain, genuine concern for the person being helped, and integrity. Human professionals are evaluated on all three. AI is evaluated almost entirely on the first. When a model consistently produces confident, well-structured answers, we

register that consistency as reliability, without any ability to verify the reasoning underneath. We have shifted from human-centered trust, grounded in relationship and accountability, toward performance-based trust measured by outputs alone.

The cycle that closes on itself. The final factor is the one I keep catching myself inside: convenience. AI is available at 3 a.m., requires no appointment, costs nothing, and carries no social friction. I saw how far this normalization has gone while searching for an AI scheduling assistant. The homepage of one chatbot greeted me with three options, side by side: an image generator, an “AI Therapy Agent,” and a writing helper. Therapy, sitting between a toy and a tool, as if it belonged there.



Invisioned.io homepage. Screenshot by the author, retrieved December 2025. The page has since changed.

Over time, this frictionless access compounds into what researchers call cognitive offloading, the outsourcing of mental tasks we would otherwise do ourselves. A 2025 study of 666 participants found that the youngest users, ages 17 to 25 (my generation), showed both the highest AI dependency and the lowest critical thinking scores.

But the finding that genuinely surprised me is who uses AI the most: people with doctorates accounted for 48.5% of heavy users, the single

largest group by education level. I expected the opposite. The data suggests higher education buffers some of the cognitive costs; doctoral holders are both the heaviest users and the best equipped to catch errors. Which raises an uncomfortable question: if even the most credentialed group needs that buffer to stay critical, what does the dependency cycle look like for the rest of us?

Because the cycle does close on itself. We use AI because it is convenient. Our skills atrophy from disuse. As dependency deepens, our capacity to critically evaluate AI's outputs weakens further. We trust it more, not because it has earned that trust, but because we have gradually lost the tools to withhold it.

Trusting in a vacuum. When I started this research, I assumed there must be laws protecting the people placing all this trust. I was shocked to find there are not, at least not in the United States. No federal legislation establishes broad regulatory authority over AI systems, and the current administration has moved to preempt state-level rules. Colorado became the first state to pass a comprehensive AI law in May 2024, but it never took effect: a federal court paused its enforcement in April 2026, and weeks later the state replaced it with a far narrower transparency law that will not apply until 2027. The EU's AI Act, in force since 2024, remains the only major framework imposing structured obligations on developers. American users remain largely unprotected, developing dependency on systems that nobody is currently obligated to make accountable.

The yes-man in the machine is agreeable by design, opaque by architecture, and indispensable by habit. Whether the trust we have placed in it reflects something AI has genuinely earned, or something we have gradually lost the capacity to question, is a problem I think we need to take seriously before the answer becomes irreversible. I know, because I started this research by questioning my own.

Next in this series: *“The black box we believe in.”* When AI fails in hospitals, courtrooms, and customer service, who is held responsible? We look at the Air Canada chatbot ruling and the growing divide between US and EU regulation.

Sources

- [1] Intelligent.com student survey, May 2023 (n=801) [Link](#)
- [2] Digital Education Council, Global AI Student Survey 2024; Latin America Survey 2026 [Link](#)
- [3] De Vynck & Merrill, Washington Post, Nov. 2025 — analysis of 47,000 ChatGPT conversations [Link](#)
- [4] Gerlich, M. (2025). AI tools in society. *Societies*, 15(1). doi:10.3390/soc15010006 [Link](#)
- [5] Mayer, Davis & Schoorman (1995). An integrative model of organizational trust. *AMR*, 20(3) [Link](#)
- [6] Sullins, J.P. (2022). Questioning trust in autonomous machines. [Link](#)
- [7] Congressional Research Service R48555 (2025); Skadden analysis of Colorado SB 24-205 [Link](#)

Disclosure: AI tools were used to assist in editing and formatting this article. All arguments, source selection, and analytical judgments are the author’s own, drawn from original seminar research conducted in Fall 2025.

Algunas entregas aparecerán también en español.